

RAVEESH RAJ GRANDHI

New York, NY | (414) 539-8603 | raveeshraj26@gmail.com | [LinkedIn](#) | [GitHub](#) | [Portfolio](#)

Data & Business Intelligence Analyst | SQL • Python • dbt • Snowflake • BigQuery | Tableau • Power BI | Experimentation & AI

WORK EXPERIENCE

Data Analyst | NYC Health + Hospitals, Woodhull Medical Center Oct 2025 – Present

- Reduced recurring report preparation time 75% (4 hours to 1) by automating data pulls, dashboard refreshes, and Excel tracking workflows across Epic, Snowflake, and internal source systems.
- Raised reporting accuracy from 88% to 94% by flagging data anomalies, tracing errors to source systems, and implementing validation checks before publication.
- Delivered daily, weekly, and monthly leadership reporting across three departments by building and maintaining 20+ Tableau, Power BI, and Excel dashboards tracking operational and productivity KPIs.
- Guided 5+ leadership performance reviews per quarter across five service lines by analyzing operational trends and variances and translating findings into clear recommendations for non-technical stakeholders.
- Ensured reporting continuity across teams by documenting data definitions, source systems, refresh cadences, and validation steps as standard process.
- Contributed to a 25% reduction in help-desk tickets during an Epic go-live by testing and validating reports and training ~50 clinical and administrative staff.

Graduate Research Analyst | Northwestern Mutual Data Science Institute Jul 2024 – May 2025

- Achieved 93.5% classification accuracy on 2,500+ clinical audio recordings by building an end-to-end Python pipeline (pandas, scikit-learn) with cross-validation and hypothesis testing — Best Poster Award, 2024.

SELECTED PROJECTS

20 tested, documented analytics projects (190+ CI tests) — github.com/Raveesh-raig

A/B Testing & Experimentation Framework | Python, scipy, statsmodels

- Quantified why experiment 'peeking' is dangerous — 24% false-positive rate for naive daily checks vs 1.2% under always-valid sequential inference (mSPRT) — by building a frequentist + Bayesian experimentation framework with power analysis, SRM gates, and 38 automated tests.

E-Commerce Analytics Platform | dbt Core, Snowflake, GitHub Actions

- Modeled 1.5M+ rows into a tested Kimball star schema (17 models, SCD2 snapshots, incremental marts) with a CI/CD pipeline building every pull request into an ephemeral schema — zero broken deploys.

Analytics RAG Agent (AI) | Python, LangChain, Claude API, DuckDB

- Enabled natural-language analytics over a governed warehouse by building an AI agent that routes questions between documentation retrieval (with cited sources) and guarded text-to-SQL — read-only execution, table allowlists, and 100% routing/retrieval accuracy on a 33-test evaluation suite.

TECHNICAL SKILLS

SQL & Modeling: Advanced SQL (window functions, CTEs, query optimization), dimensional/star-schema modeling, dbt Core, ETL/ELT, data quality testing, Snowflake, BigQuery, DuckDB, PostgreSQL

Analytics & Statistics: Python (pandas, scikit-learn, statsmodels), A/B testing & experiment design, causal inference, cohort & funnel analysis, forecasting, KPI definition

BI & AI: Tableau, Power BI (DAX), Looker Studio, Excel (Power Query), Streamlit; LLM integration with evaluation harnesses (Claude, Gemini), RAG, prompt guardrails; Git/GitHub, CI/CD, Jira

EDUCATION & CERTIFICATIONS

M.S. Healthcare Informatics — University of Wisconsin–Milwaukee | Chancellor's Graduate Student Award

Google Advanced Data Analytics Certificate • Best Poster Award, NM Data Science Institute (2024)